

Recurrent Neural Networks



Ha-Jin Yu, Dept. of Computer Science, University of Seoul

서울시립대학교 컴퓨터과학부 유하진

2019


HJYU@UOS.AC.KR



서울시립대학교
UNIVERSITY OF SEOUL

Sequence Modeling

- Recurrent neural networks or RNNs (Rumelhart et al., 1986a)
- Sequential data $x^{(1)}, \dots, x^{(\tau)}$ 처리 (다양한 길이)
 - Speech, text, video, handwriting, ...



RNN applications

- Unconstrained handwriting recognition (graves *et al.*, 2009),
- Speech recognition (graves *et al.*, 2013; graves and jaitly, 2014),
- Handwriting generation (graves, 2013),
- Machine translation (sutskever *et al.*, 2014),
- Image captioning (kiros *et al.*, 2014b; vinyals *et al.*, 2014b; xu *et al.*, 2015)

Generating hand-writing using RNN

from his travels it might have been
from his travels it might have been
from his travels it might have been
from his travels it might have been

university of seoul
university of seoul
university of seoul


Graves, Alex. "Generating sequences with recurrent neural networks." *arXiv preprint arXiv:1308.0850* (2013).



RNN (Recurrent Neural Network)

- Feedforward : $\mathbf{h}(t) = f(\mathbf{x}(t); \theta)$

- RNN : $\mathbf{h}(t) = f(\mathbf{h}(t-1), \mathbf{x}(t); \theta)$



Simple RNN

- Feedforward : $\mathbf{h}(t) = f(\mathbf{x}(t); \theta)$
- RNN : $\mathbf{h}(t) = f(\mathbf{h}(t-1), \mathbf{x}(t); \theta)$

$$\mathbf{h}(t) = \tanh(\mathbf{W}_{\text{hh}} \mathbf{h}(t-1) + \mathbf{W}_{\text{xh}} \mathbf{x}(t))$$
$$\mathbf{y}(t) = \mathbf{W}_{\text{hy}} \mathbf{h}(t)$$

$$\begin{aligned}\mathbf{h}(t) &= \tanh(W_{hh}h_{t-1} + W_{xh}x_t) \\ &= \tanh \left(\begin{pmatrix} W_{hh} & W_{hx} \end{pmatrix} \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right) \\ &= \tanh \left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right)\end{aligned}$$


- Also called as “**Vanilla RNN**” or “Elman RNN” after Prof. Jeffrey Elman



RNN: Computational Graph


h_t : time t에서 hidden node들의 값

x_t : time t에서 input x 값




<http://cs231n.stanford.edu/>

RNN: Computational Graph



<http://cs231n.stanford.edu/>


RNN: Computational Graph



<http://cs231n.stanford.edu/>


RNN: Computational Graph

Re-use the same weight matrix at every time-step




<http://cs231n.stanford.edu/>

RNN: Computational Graph: Many to Many



<http://cs231n.stanford.edu/>


RNN: Computational Graph: Many to Many



e.g. **Machine Translation**
seq of words → seq of words

<http://cs231n.stanford.edu/>

Unfolding



$t = 13$

$t = 14$

$t = 15$


$t = 16$

- 매 시간 별 상태를 보여주는 그림

RNN의 기본적인 구조


loop가 존재하여,
이전 단계의 정보가
현재 단계에 반영되도록 함

unfolding




x_t : 시간 t 에서의 입력
 h_t : 시간 t 에서의 출력

RNN의 unfolding



RNN의 기본적인 구조 (Many to Many)

- Recurrent connections between hidden units




<http://www.deeplearningbook.org>




RNN 구조 (Many to One)

- A single output at the end of the sequence.



<http://www.deeplearningbook.org>


RNN: Computational Graph: Many to One



e.g. **speech recognition**
Speech waveform → word

<http://cs231n.stanford.edu/>

RNN: Computational Graph: One to Many



e.g. **Image Captioning**
image → sequence of words

<http://cs231n.stanford.edu/>

영상 주석 생성

- <http://deeplearning.cs.toronto.edu/i2t>



a man wearing a blue shirt with his arms on the grass.
a man holding a frisbee bat in front of a green field.
a man throwing a frisbee in a green field.
a boy playing ball with a disc in a field.
a young man playing in the grass with a green ball.



a red car on the side of the road in the small race.
a truck driving uphill on the side of the road.
a person driving a truck on the road.
a small car driving down a dirt and water.
a truck in a field of car is pulled up to the back.




a group of birds standing next to each other.
a group of ducks that are standing in a row.
a group of ducks that are standing on each other.
a group of sheep next to each other on sand.
a group of small birds is standing in the grass.



a kite flying over the ocean on a sunny day.
a person flying over the ocean on a sunny day.
a person flying over the ocean on a cloudy day.
a kite on the beach on the water in the sky.
a large flying over the water and rocks.



Encoder-Decoder Sequence-to-Sequence Architectures




RNN의 학습 : *BPTT*

- Back-propagation algorithm 그대로 적용
- *Back-propagation through time*
 - The back-propagation algorithm applied to the unrolled graph
- Forward pass 과정의 States 는 backward pass 계산을 위해서 저장됨
- Memory cost is $O(\tau)$.



Multilayer RNNs



$$h_t^l = \tanh W^l \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix}$$



$$h \in \mathbb{R}^n$$

$$W^l \ [n \times 2n]$$

<http://cs231n.stanford.edu/>


The Challenge of Long-Term Dependencies

- 예: RNN을 이용해 문장의 마지막 단어를 추측
 - “The clouds are in the ()”
 - 추측 대상인 ‘sky’와 추측에 필요한 ‘clouds’ 간의 간격이 작음
 - RNN으로 추측 가능
 - “I grew up in France ... (중략) ... I speak fluent ()”
 - ‘speak’을 통해 마지막 단어는 언어 이름일 것임을 추측 가능
 - 어떤 언어인지까지 추측하고 싶다면?
 - 문장 맨 앞부분에 ‘France’가 나왔다는 정보가 추가로 필요
 - **RNN으로는 추측이 어려움**
 - » ‘France’와 ‘French’ 간 간격이 크기 때문 (long-term dependency)




Long-term dependencies 처리 방법

- Multiple time scales 을 처리하는 model
 - Fine-grained time scales - small details
 - Coarse time scales – 먼 과거에서 현재로 정보 전달
- 해결 방법 들
 - Skip connections
 - Leaky units :
Hidden units with linear self-connections
 - Removing Connections:
removing length-one
connections and replacing
them with longer connections.




Leaky Units

- Integrate signals with different time constants
- $M^{(t)} \leftarrow \alpha\mu^{(t)} + (1 - \alpha)v^{(t)}$
- $M^{(t)}$: running average of some value $v^{(t)}$
- Linear self-connection from $\mu(t-1)$ to $\mu(t)$.
- α parameter : when $\alpha = 1$: 오랫동안 기억
 $\alpha = 0$: 빨리 잊어버림
- 얼마나 오래 전 정보까지 기억할지를 실수 α 로 결정
- α 는 상수이거나 또는 학습되거나




Gated RNNs

- Gating units : controls the flow of information
- Gated RNNs : leaky unit의 일반화.
Connection weights 가 시간에 따라 변할 수 있음.
- Leaky units : 정보를 오랫동안 축적
- 정보가 사용되고 나면 잊어버릴 (forget) 필요가 있음.
- 필요한 구간의 정보를 처리하고 나면 α 를 0으로 setting 하도록 학습
- Gated : self-loop의 weight를 다른 hidden unit 이 제어하도록하여 입력에 따라 시간 scale 을 dynamically 변경
- Forget gate
 - 이전 단계의 정보 중, 현재 단계부터는 더 이상 필요하지 않은 정보를 제거



Long Short Term Memory (LSTM)




LSTM memory cell

<http://deeplearning.net/tutorial/lstm.html>

LSTM recurrent network “cell.”

<http://www.deeplearningbook.org>



Forget gate unit

- The **Self-loop weight** is controlled by a *forget gate* unit $f_i^{(t)}$ (for time step t and cell i)
- A **sigmoid** unit to obtain a gating value between 0 and 1

$$f_i^{(t)} = \sigma \left(b_i^f + \sum_j U_{i,j}^f x_j^{(t)} + \sum_j W_{i,j}^f h_j^{(t-1)} \right)$$

- $x(t)$ = current input vector
- $h(t)$ = the current hidden layer vector, containing the outputs of all the LSTM cells
- b^f = biases
- U^f = input weights
- W^f = recurrent weights for the forget gates.

LSTM cell internal state $s_i^{(t)}$

$$s_i^{(t)} = f_i^{(t)} s_i^{(t-1)} + g_i^{(t)} \tanh \left(b_i + \sum_j U_{i,j} x_j^{(t)} + \sum_j W_{i,j} h_j^{(t-1)} \right)$$

- $f_i^{(t)}$:conditional self-loop weight
- b :biases,
- U :input weights
- W :recurrent weights into the LSTM cell.

The *external input gate* unit $g_i^{(t)}$

- A Sigmoid unit to obtain a gating value between 0 and 1

$$g_i^{(t)} = \sigma \left(b_i^g + \sum_j U_{i,j}^g x_j^{(t)} + \sum_j W_{i,j}^g h_j^{(t-1)} \right)$$



Output gate

- The output $h_i^{(t)}$ of the LSTM cell can also be **shut off**, via the *output gate* $q_i^{(t)}$
- also uses a **sigmoid** unit for gating:

$$q_i^{(t)} = \sigma \left(b_i^o + \sum_j U_{i,j}^o x_j^{(t)} + \sum_j W_{i,j}^o h_j^{(t-1)} \right)$$

$$h_i^{(t)} = \tanh \left(s_i^{(t)} \right) q_i^{(t)}$$

- parameters
 - b^o :biases,
 - U^o :input weights,
 - W^o :recurrent weights

Long Short Term Memory (LSTM)

- 모든 weights를 W 하나로 표현하면

Vanilla RNN

$$h_t = \tanh \left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right)$$

stack

LSTM

$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$


state $c_t = f \odot c_{t-1} + i \odot g$

$$h_t = o \odot \tanh(c_t)$$

Hochreiter and Schmidhuber, “Long Short Term Memory”, Neural Computation, 1997
<http://cs231n.stanford.edu/>



Long Short Term Memory (LSTM)



i: Input gate, whether to write to cell
f: Forget gate, Whether to erase cell
o: Output gate, How much to reveal cell


$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$c_t = f \odot c_{t-1} + i \odot g$$
$$h_t = o \odot \tanh(c_t)$$

<http://cs231n.stanford.edu/>



LSTM




$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$c_t = f \odot c_{t-1} + i \odot g$$


$$h_t = o \odot \tanh(c_t)$$

<http://cs231n.stanford.edu/>

Vanilla RNN Gradient Flow



$$\begin{aligned} h_t &= \tanh(W_{hh}h_{t-1} + W_{xh}x_t) \\ &= \tanh \left(\begin{pmatrix} W_{hh} & W_{hx} \end{pmatrix} \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right) \\ &= \tanh \left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right) \end{aligned}$$





Backpropagation 할 때 W 가
반복해서 곱해짐 →
Exploding gradients
Vanishing gradients

<http://cs231n.stanford.edu/>


Bengio et al, "Learning long-term dependencies with gradient descent is difficult", IEEE Transactions on Neural Networks, 1994

Pascanu et al, "On the difficulty of training recurrent neural networks", ICML 2013


Recall: gradient calculation



- 곱셈보다 덧셈이 유리함.



LSTM: Gradient Flow



$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$


$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$

<http://cs231n.stanford.edu/>

LSTM: Gradient Flow

gradient 계산할 때 W 가 곱해지지 않음




f 값은 일정하지 않음.

→ Vanishing gradients 문제 해결

<http://cs231n.stanford.edu/>

Gated Recurrent Unit (GRU)

- Two gates: reset gate r and update gate z
- No Output gate





$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Residual Net과 비슷한 원리



- Gradient Flow 를 위한 통로를 만들어 깊은 층 구현 가능